

Jeff Caldwell

CS 3311

Assignment 01

September 13, 2021

Comparison of Sequential and Random File Reading Operations in C

In order to compare sequential and random file read operations, three functions were created using the C programming language. C was chosen for this comparison in order to avoid any compiler optimization that might skew the results.

The first function of the three, `writeRecord`, creates a binary file and writes 100 thousand simple data structures composed of an incrementing `ID` member of type `integer` and a `Name` member of type `char` with a size of `1020` bytes.

```
struct MyRecord {  
    int ID;  
  
    char Name[1020];  
}
```

Record to be written to binary file

```
void writeRecord() {  
    FILE *fp;  
  
    fp = fopen("data.bin", "wb");  
  
    struct MyRecord rec;  
  
    if(fp == NULL) {  
        printf("Error opening data.bin\n");  
  
        exit(1);  
    } else {  
  
        for(int i = 1; i <= 100000; i++) {  
            rec.ID = i;  
            strcpy(rec.Name, "ABCDE");  
        }  
    }  
}
```

```

        fwrite(&rec, sizeof(struct MyRecord), 1, fp);
    }
}

fclose(fp);
}

```

Function to write 100,000 records to a binary file

After writing the binary file, two additional functions were called — `readAllData`, and `readDataWithRandom`. These functions, respectively, read the data from the written binary file in sequence, and read the data from the file by selecting and reading a record of 1024 bytes in size from a random position in the file.

```

void readAllData() {
    FILE *fp;
    struct MyRecord rec;

    fp = fopen("data.bin", "rb");

    if(fp == NULL) {
        printf("Error opening data.bin\n");

        exit(1);
    } else {

        for(int i = 1; i <= 100000; i++) {
            fread(&rec, sizeof(struct MyRecord), 1, fp);
        }

        fclose(fp);
    }
}

```

Function to read all 100,000 records from the binary file sequentially

```

void readDataWithRandom() {
    FILE *fp;

    struct MyRecord rec;

    fp = fopen("data.bin", "rb");

    if(fp == NULL) {
        printf("Error opening data.bin.");

        exit(1);
    }
}

```

```

} else {
    for(int i = 0; i < 100000; i++) {
        int rando = getRandom(1, 100000);

        fseek(fp, sizeof(struct MyRecord) * rando, SEEK_SET);

        fread(&rec, sizeof(struct MyRecord), 1, fp);
    }
}

fclose(fp);
}

```

Function to read a record from a random position in the binary file 100,000 times

In order to measure and compare the performance of each function, each was called 50 times and the amount of time for each function was measured using `clock()`. The data from each full iteration of the report was then stored.

```

void generateReadWriteReports(int iterations) {

    struct ReportData allReports[iterations];

    for(int i = 0; i < iterations; i++) {
        clock_t
            w_start,
            w_end,
            r_start,
            r_end;

        double w_time, r_time;

        struct ReportData report;

        writeRecord();

        w_start = clock();
        readAllData();
        w_end = clock();
        w_time = ((double) w_end - w_start) / CLOCKS_PER_SEC;
        report.read = w_time;

        r_start = clock();
        readDataWithRandom();
        r_end = clock();
        r_time = ((double) r_end - r_start) / CLOCKS_PER_SEC;
        report.random = r_time;
    }
}

```

```

    if(r_time > w_time) {
        report.difference = r_time - w_time;
        report.longer = "Random";
    } else {
        report.difference = w_time - r_time;
        report.longer = "Nonrandom";
    }

    allReports[i] = report;
}

printReport(iterations, allReports);
}

```

Function to execute and time the separate sequential and random read operations

Following all 50 iterations of `generateReadWriteReport`, the results were printed to the terminal. The printed report included the time to read the file sequentially, or `Read Time`, the time to read from a random position, or `Random Time`, the difference between the two execution times: `Difference`, and the name of which function took longer to execute: `Longer`.

Iteration	Read Time	Random Time	Difference	Longer
1	0.028697s	0.127618s	0.098921s	Random
2	0.028136s	0.130952s	0.102816s	Random
3	0.028668s	0.126355s	0.097687s	Random
4	0.029613s	0.129292s	0.099679s	Random
5	0.028335s	0.126712s	0.098377s	Random
6	0.028406s	0.132453s	0.104047s	Random
7	0.028286s	0.125564s	0.097278s	Random
8	0.028358s	0.125410s	0.097052s	Random
9	0.028039s	0.127805s	0.099766s	Random
10	0.028134s	0.129565s	0.101431s	Random
11	0.028356s	0.133298s	0.104942s	Random
12	0.028620s	0.129533s	0.100913s	Random
13	0.029001s	0.133559s	0.104558s	Random
14	0.028436s	0.122594s	0.094158s	Random
15	0.029554s	0.126746s	0.097192s	Random
16	0.028829s	0.129995s	0.101166s	Random
17	0.027992s	0.125700s	0.097708s	Random
18	0.028860s	0.128784s	0.099924s	Random
19	0.028627s	0.130293s	0.101666s	Random
20	0.030559s	0.126866s	0.096307s	Random
21	0.028414s	0.125810s	0.097396s	Random
22	0.029059s	0.131566s	0.102507s	Random
23	0.028289s	0.141398s	0.113109s	Random
24	0.031719s	0.157298s	0.125579s	Random

25	0.036970s	0.182742s	0.145772s	Random
26	0.042999s	0.132257s	0.089258s	Random
27	0.036828s	0.161207s	0.124379s	Random
28	0.034316s	0.154184s	0.119868s	Random
29	0.033394s	0.153454s	0.120060s	Random
30	0.030945s	0.150024s	0.119079s	Random
31	0.027844s	0.128146s	0.100302s	Random
32	0.028071s	0.131426s	0.103355s	Random
33	0.028594s	0.127795s	0.099201s	Random
34	0.027230s	0.126656s	0.099426s	Random
35	0.027371s	0.123966s	0.096595s	Random
36	0.030658s	0.128724s	0.098066s	Random
37	0.028103s	0.125251s	0.097148s	Random
38	0.028259s	0.128333s	0.100074s	Random
39	0.028535s	0.132456s	0.103921s	Random
40	0.028215s	0.129247s	0.101032s	Random
41	0.028203s	0.128128s	0.099925s	Random
42	0.028103s	0.126843s	0.098740s	Random
43	0.028434s	0.126222s	0.097788s	Random
44	0.028093s	0.133815s	0.105722s	Random
45	0.029312s	0.126838s	0.097526s	Random
46	0.028267s	0.131841s	0.103574s	Random
47	0.028757s	0.125956s	0.097199s	Random
48	0.029394s	0.125578s	0.096184s	Random
49	0.028082s	0.129016s	0.100934s	Random
50	0.027908s	0.126329s	0.098421s	Random

Average:	0.029477s	0.132432s	0.102955s	

Terminal output of 50 executions of both sequential and random read operations

On a computer with a 6 core, 2.3GHz AMD processor, and 32 GiB of RAM, reading from a 1TB M.2 solid-state drive, the random-selection function took the longest time to execute in every case.

After running this experiment multiple times, it is clear that random I/O takes longer to execute than sequential I/O. It's possible that the additional processing time required to calculate the random location within the file, in addition to the time required for the solid-state drive to process the read operations is the reason for the disparity.

While the small amount of data collected from this experiment cannot be conclusive, it does indicate that enforcing sequential read/write operations wherever possible is the most efficient way to perform disk I/O.

Appendix

main.c

```
// Jeff Caldwell
// CS 3311
```

```

// Assignment 01
// September 13, 2021

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>

struct MyRecord {
    int ID;

    char Name[1020];
};

struct ReportData {
    double read;
    double random;
    double difference;
    char* longer;
};

void writeRecord();
void readAllData();
void readDataWithRandom();
void generateReadWriteReports(int iterations);
int getRandom(int min, int max);
void printReport(int iterations, struct ReportData records[]);

int main() {

    printf("Reading data...\n");
    generateReadWriteReports(50);

    return 0;
}

void writeRecord() {
    FILE *fp;

    fp = fopen("data.bin", "wb");

    struct MyRecord rec;

    if(fp == NULL) {
        printf("Error opening data.bin\n");

        exit(1);
    } else {

```

```

    for(int i = 1; i <= 100000; i++) {
        rec.ID = i;
        strcpy(rec.Name, "ABCDE");

        fwrite(&rec, sizeof(struct MyRecord), 1, fp);
    }
}

fclose(fp);
}

void readAllData() {
    FILE *fp;
    struct MyRecord rec;

    fp = fopen("data.bin", "rb");

    if(fp == NULL) {
        printf("Error opening data.bin\n");

        exit(1);
    } else {

        for(int i = 1; i <= 100000; i++) {
            fread(&rec, sizeof(struct MyRecord), 1, fp);
        }

        fclose(fp);
    }
}

void readDataWithRandom() {
    FILE *fp;

    struct MyRecord rec;

    fp = fopen("data.bin", "rb");

    if(fp == NULL) {
        printf("Error opening data.bin.");

        exit(1);
    } else {
        for(int i = 0; i < 100000; i++) {
            int rando = getRandom(1, 100000);

            fseek(fp, sizeof(struct MyRecord) * rando, SEEK_SET);

            fread(&rec, sizeof(struct MyRecord), 1, fp);
            // printf("{ID: %d, Name: %s}", rec.ID, rec.Name);

```

```

    }
}

fclose(fp);
}

void generateReadWriteReports(int iterations) {

    struct ReportData allReports[iterations];

    for(int i = 0; i < iterations; i++) {
        clock_t
            w_start,
            w_end,
            r_start,
            r_end;

        double w_time, r_time;

        struct ReportData report;

        writeRecord();

        w_start = clock();
        readAllData();
        w_end = clock();
        w_time = ((double) w_end - w_start) / CLOCKS_PER_SEC;
        report.read = w_time;

        r_start = clock();
        readDataWithRandom();
        r_end = clock();
        r_time = ((double) r_end - r_start) / CLOCKS_PER_SEC;
        report.random = r_time;

        if(r_time > w_time) {
            report.difference = r_time - w_time;
            report.longer = "Random";
        } else {
            report.difference = w_time - r_time;
            report.longer = "Nonrandom";
        }

        allReports[i] = report;
    }

    printReport(iterations, allReports);
}

void printReport(int iterations, struct ReportData records[]) {

```



```

double
    totalRead,
    averageRead,
    totalRandom,
    averageRandom,
    totalDifference,
    averageDifference;

printf("-----\n");
printf("Iteration\tRead Time\tRandom Time\tDifference\tLonger\n");
printf("-----\n");
for(int i = 0; i < iterations; i++) {
    totalRead = totalRead + records[i].read;
    totalRandom = totalRandom + records[i].random;
    totalDifference = totalDifference + records[i].difference;

    printf(
        "%d\t\t%s\t%s\t%s\t%s\n",
        i + 1,
        records[i].read,
        records[i].random,
        records[i].difference,
        records[i].longer
    );
}
printf("-----\n");

averageRead = totalRead / iterations;
averageRandom = totalRandom / iterations;
averageDifference = totalDifference / iterations;

printf("Average:\t%s\t%s\t%s\n", averageRead, averageRandom,
averageDifference);
printf("-----\n");
}

int getRandom(int min, int max) {
    return (rand() % (max - min + 1) + min);
}

```